The Safety Standoff

Anthropic was punished for demanding AI safety restrictions the Pentagon quietly accepted from OpenAI days later. This is the story of that contradiction - and why it matters.

KEVIN LU · NORTHEASTERN UNIVERSITY · APRIL 2026

INTRODUCTION

One Company Punished, Another Rewarded - for the Same Thing

In late February 2026, the U.S. Department of Defense and Anthropic, the AI company behind the chatbot Claude, had a very public falling out. The core issue was simple: Anthropic did not want Claude used for mass domestic surveillance or in fully autonomous weapons, and the Pentagon wanted the right to use it for "all lawful purposes" with no restrictions.

When Anthropic refused to back down, the government canceled its $200 million contract, designated the company a "supply chain risk" - a label normally reserved for foreign adversaries like Russia and China - and ordered every federal agency to stop using Anthropic's products. Then, just hours later, OpenAI announced it had signed its own classified-network deal with the Pentagon that included essentially the same prohibitions Anthropic was punished for demanding.

The dispute has since escalated into active litigation, with Anthropic filing two federal lawsuits alleging that the government retaliated against the company for its speech on AI safety. This story touches on civil liberties, the legal limits of executive power, the role of AI in warfare, and whether a private company can set ethical conditions on technology it sells to the military.

TIMELINE

Four Weeks That Shook the AI Industry

What started as a contract dispute between one company and the Pentagon quickly became the most significant confrontation between the tech industry and the U.S. government since the Snowden revelations. Here is how it unfolded.

FEB 24

Hegseth Delivers Ultimatum

Defense Secretary calls Amodei to Pentagon; demands all restrictions dropped by Friday 5:01 PM.

READ MORE

FEB 25

Lawfare Legal Analysis Published

Rozenshtein publishes detailed analysis of DPA's applicability to AI software disputes.

READ MORE

FEB 26

Amodei's Public Statement

Anthropic CEO publicly lays out two red lines: no mass surveillance, no autonomous weapons.

READ MORE

FEB 27

Contract Canceled, Anthropic Blacklisted

Pentagon cancels $200M contract; designates Anthropic a "supply chain risk"; all agencies ordered to stop using Anthropic products.

READ MORE

FEB 28

OpenAI Signs Pentagon Deal

OpenAI announces classified-network contract with the same core prohibitions Anthropic was punished for.

READ MORE

FEB 27

"We Will Not Be Divided" Letter

300+ Google and 60+ OpenAI employees sign open letter backing Anthropic's red lines.

READ MORE

MAR 5

Anthropic Responds to Blacklisting

Amodei clarifies the supply chain label's scope and announces intent to challenge it in court.

READ MORE

MAR 9

Two Federal Lawsuits Filed

Anthropic sues in San Francisco and D.C. Circuit, alleging First Amendment retaliation.

READ MORE

MAR 16–26

Court Rules: "Orwellian" Retaliation

After amicus briefs and a March 24 hearing, Judge Rita F. Lin issues a 43-page ruling granting Anthropic a preliminary injunction.

READ MORE

FIGURE 1 - KEY EVENTS IN THE ANTHROPIC–PENTAGON DISPUTE, FEB 24 – MAR 24, 2026

THE CONFRONTATION

An Ultimatum From the Pentagon

On February 24, Defense Secretary Pete Hegseth called Anthropic CEO Dario Amodei to the Pentagon and delivered an ultimatum: drop all restrictions on how the military can use Claude by Friday at 5:01 PM, or face consequences that could include being forced to cooperate under the Defense Production Act - a Korean War-era statute normally used to compel production of physical goods like steel and ammunition.

Amodei refused. In a public statement two days later, he laid out Anthropic's two red lines and the reasoning behind each.

> *On surveillance: current law has not caught up with AI's capabilities. The government can already buy Americans' location data, browsing history, and financial records from data brokers without a warrant. AI makes it possible to take all of that scattered information and automatically assemble it into a detailed profile of any individual at massive scale.*
>
> - Dario Amodei, CEO of Anthropic, February 26, 2026

On autonomous weapons, Amodei's argument was more technical: frontier AI models are simply not reliable enough to make life-or-death targeting decisions without a human in the loop. He also pointed out a contradiction in the

government's position - one arm calling Anthropic a security risk while another insisted Claude was essential to national security.

THE CONTRADICTION

Same Safeguards, Different Outcomes

Hours after Anthropic was blacklisted, OpenAI published a blog post describing its new Pentagon contract. The deal included three explicit prohibitions: no mass domestic surveillance, no directing autonomous weapons, and no social credit-type systems. OpenAI also said it asked the Pentagon to offer the same terms to all AI labs, including Anthropic.

The overlap is striking. Compare the two companies' positions side by side.

Anthropic's Red Lines	OpenAI's Prohibitions
REJECTED BY PENTAGON	ACCEPTED BY PENTAGON
No mass domestic surveillance of American citizens	No mass domestic surveillance
No deployment in fully autonomous weapons without human oversight	No directing autonomous weapons
	No social credit-type scoring systems

Reveal the Overlap

FIGURE 2 - ANTHROPIC'S TWO RED LINES VS. OPENAI'S THREE PROHIBITIONS. THE CORE RESTRICTIONS ARE IDENTICAL.

If the Pentagon accepted these safeguards from OpenAI, then the dispute with Anthropic was never really about the safeguards themselves. It was about who

gets to set conditions. Sam Altman later admitted the OpenAI deal was rushed and acknowledged it looked "opportunistic and sloppy." The contract language actually went through revisions after legal experts found that the initial wording left loopholes around commercially acquired data.

If the Pentagon accepted these safeguards from OpenAI, then the dispute with Anthropic was never really about the safeguards themselves. It was about who gets to set conditions.

THE LEGAL BATTLE

From Contract Dispute to Constitutional Fight

The legal situation is messier than either side would probably like to admit. The Defense Production Act is a Korean War-era statute that was built for physical goods like steel and aluminum, not disputes about software safety features.

> *There is a useful distinction between two different things the government might demand: dropping the contractual usage restrictions on Claude - legally contested but arguably within the DPA's reach - versus forcing Anthropic to retrain the model to remove its safety behavior, which starts to look like demanding a product the company doesn't make.*
>
> - Alan Z. Rozenshtein, law professor and former DOJ attorney, writing in Lawfare

Rozenshtein's biggest takeaway is that this entire conflict exists because Congress has never passed legislation setting rules for how the military can use AI. So the

question of whether the Pentagon can use Claude for mass surveillance comes down to whoever happens to be running the Department of Defense at the time.

Internal Pentagon records later revealed that Anthropic was designated a supply chain risk not because of any legitimate security concern, but because of its "hostile manner through the press" - in other words, for speaking publicly about AI safety.

On March 9, Anthropic took the fight to court, filing two federal lawsuits: one in San Francisco and another in the D.C. Circuit Court of Appeals. The suits argue the supply chain risk label was illegal retaliation for Anthropic's protected speech about AI safety, not a legitimate national security determination.

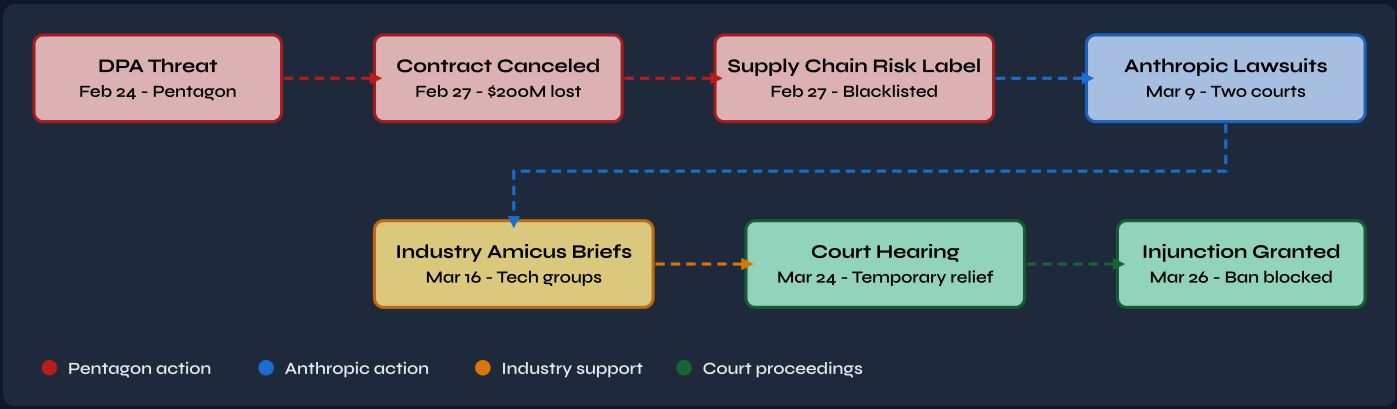

FIGURE 3 - LEGAL PATHWAY FROM PENTAGON THREAT TO PENDING COURT RULING

> *If the government can blacklist a company and claim it's a security risk over a contract disagreement, the procurement system becomes contingent on political favor rather than the rule of law.*
>
> - Tech industry amicus brief filed March 16, 2026

Major tech industry trade groups filed amicus briefs making the point that the stakes go well beyond Anthropic. The Pentagon's Undersecretary Emil Michael said the department is looking to fully replace Anthropic, and a court hearing on whether to grant temporary relief was scheduled for March 24.

On March 26, U.S. District Judge Rita F. Lin issued a stinging 43-page ruling granting Anthropic a preliminary injunction, blocking the government from enforcing its ban and the supply chain risk designation.

> *Nothing in the governing statute supports the Orwellian notion that an American company may be branded a potential adversary and saboteur of the U.S. for expressing disagreement with the government.*
>
> - Judge Rita F. Lin, U.S. District Court, Northern District of California

INDUSTRY SOLIDARITY

"We Will Not Be Divided"

One factor that complicated the Pentagon's strategy was the AI workforce. More than 300 Google employees and over 60 OpenAI employees signed an open letter titled "We Will Not Be Divided," which called on their companies to hold the same red lines as Anthropic. The letter argued that the Pentagon was deliberately trying to split the industry - making each company afraid the others would cave.

Google DeepMind Chief Scientist Jeff Dean posted on X that mass surveillance violates the Fourth Amendment. Sam Altman said publicly that he didn't think the Defense Production Act should be used against AI companies. The concerns Anthropic raised were not just one company's position - they were shared across the AI workforce.

But how durable is this solidarity? As AI policy expert Hamid Ekbia of Syracuse University points out, in a tighter job market employees may not feel as safe speaking out, which means this kind of cross-company pressure could be hard to sustain. He calls the AI industry's competitive environment a "moral economy of the jungle" where the most aggressive companies tend to win.

INTERNATIONAL CONTEXT

The World Already Agrees

The question of whether AI should be allowed to autonomously select and kill targets is not something only Anthropic cares about. In December 2024, the UN General Assembly voted **166 to 3** in favor of a resolution supporting a framework to prohibit or regulate lethal autonomous weapons systems. The UN Secretary-General and the International Committee of the Red Cross have called for a binding treaty on autonomous weapons by the end of 2026.

The U.S. abstained. The three votes against came from Russia, Belarus, and North Korea. The visual below represents every vote cast. Each square is a nation. The green is overwhelming. The red is almost invisible.

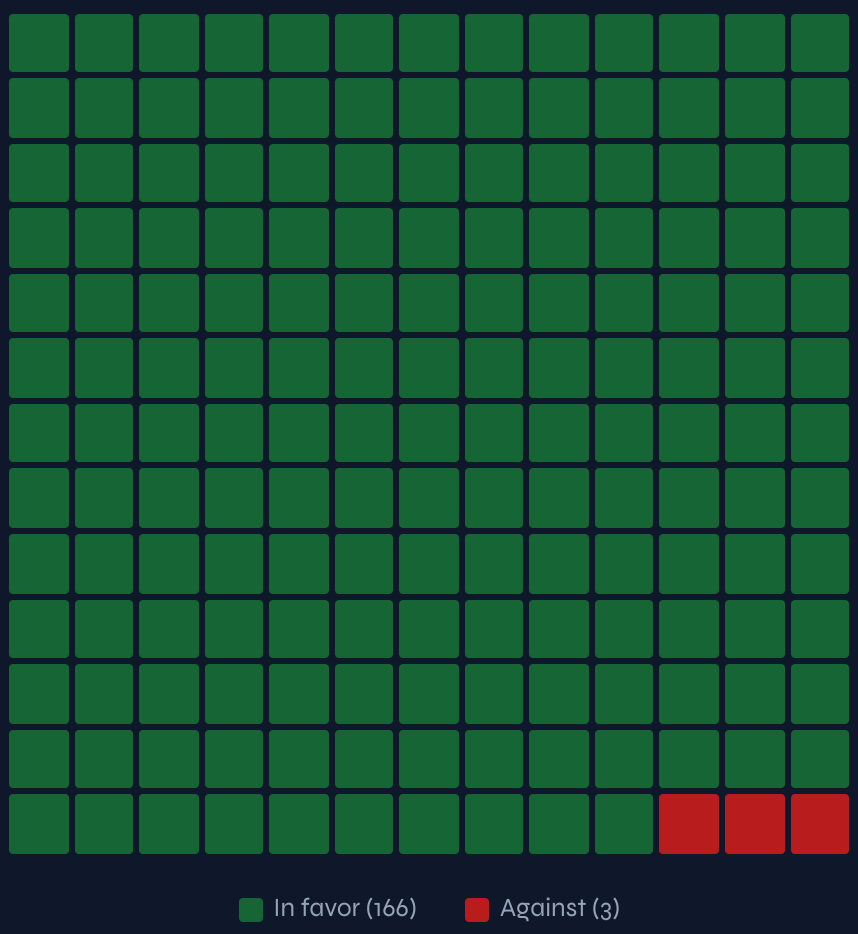

FIGURE 4 - UN GENERAL ASSEMBLY VOTE ON LETHAL AUTONOMOUS WEAPONS FRAMEWORK, DECEMBER 2024. 166 NATIONS VOTED IN FAVOR; RUSSIA, BELARUS, AND NORTH KOREA VOTED AGAINST. THE U.S. ABSTAINED.

The U.S. has resisted a preemptive ban, but the overwhelming international support shows that Anthropic's insistence on human-in-the-loop requirements aligns with a broad global consensus - not some niche concern. If the U.S. participates in and eventually ratifies such a treaty, it would create a durable legal floor for the kind of restrictions Anthropic was advocating.

SOLUTIONS

What Needs to Change

1. Federal Legislation

The fix that keeps coming up across every expert analysis is federal legislation, and it is hard to argue against. If Congress had set clear rules on autonomous weapons and surveillance, Anthropic would probably have been fine selling to the military, and none of this would have happened. Right now, the question of whether the Pentagon uses AI for mass surveillance depends entirely on who is running the Department of Defense at any given moment. Congress could pass legislation requiring judicial authorization for AI-driven domestic surveillance and mandating human oversight in any AI system involved in lethal targeting decisions. The Pentagon itself publicly said it was not interested in either of those use cases - writing them into law would take the pressure off individual companies.

2. Industry Standards

Legislation alone is not enough if individual companies can still be picked off one at a time. An industry-wide framework, modeled on the kind of professional standards that govern medicine or law, could make it much harder for the government to single out and punish one company for holding a line the entire industry agrees on. The challenge is that the competitive structure of the AI industry actively works against this kind of coordination - but the "We Will Not

Be Divided" letter is evidence that at least the workforce wants to move in that direction.

3. International Law

The UN vote of 166-3 and the Secretary-General's call for a binding treaty by end of 2026 represent the strongest possible international mandate. If ratified, such a treaty would create protections far harder to undo than any single company's contract language, signaling to the defense industry that human oversight in weapons systems is not a negotiable preference but a legal requirement.

CONCLUSION

Bigger Than One Company

This dispute is bigger than one contract or one company. It is about who gets to decide the limits on how powerful AI systems are used by the government - and what happens when a company tries to draw those lines on its own. Anthropic took a massive financial hit because its leadership decided that the risks of unrestricted military AI were not acceptable. The company was then punished for that stance while a competitor got essentially the same protections by being less confrontational about it.

The safety of these systems currently depends on the individual courage of whoever happens to be running an AI company at the time. That is not a sustainable model.

Congress needs to legislate clear boundaries. The AI industry needs collective commitments it cannot easily abandon under competitive pressure. The courts need to establish that the government cannot retaliate against companies for exercising their right to speak on AI safety. Until those structural changes are in place, the next version of this crisis is just a matter of time.

SOURCES

References

1. Allyn, B. (2026, March 9). Anthropic sues the Trump administration over 'supply chain risk' label. *NPR*. npr.org
2. Amodei, D. (2026, February 26). Statement from Dario Amodei on our discussions with the Department of War. *Anthropic*. anthropic.com
3. Anthropic. (2026, March 5). Where we stand on the Department of War. *Anthropic Blog*. anthropic.com
4. Bethe, R., & Iyer, K. (2026, February 24). Pentagon threatens to make Anthropic a pariah if it refuses to drop AI guardrails. *CNN*. cnn.com
5. Fried, I. (2026, March 16). Tech firms back Anthropic in Pentagon lawsuit. *Axios*. axios.com
6. Munoz, C. (2026, March 13). Anthropic-Pentagon dispute reveals limits of AI self-regulation, expert says. *Syracuse University Today*. news.syr.edu
7. OpenAI. (2026, February 28). Our agreement with the Department of War. *OpenAI Blog*. openai.com
8. Perrin, B. (2025, January 24). Lethal autonomous weapons systems and international law. *ASIL Insights, 29*(1). asil.org
9. Rozenshtein, A. Z. (2026, February 25). What the Defense Production Act can and can't do to Anthropic. *Lawfare*. lawfaremedia.org

10. Silberling, A. (2026, February 27). Employees at Google and OpenAI support Anthropic's Pentagon stand in open letter. *TechCrunch*. techcrunch.com

An editorial project for Advanced Writing · Northeastern University · 2026